

Parallel computations with R

Søren Højsgaard

Aalborg University, Denmark

October 15, 2012

Contents

1	Introduction	3
2	Parametric bootstrap	4
2.1	A logistic regression	4
2.2	Parametric bootstrap	8
2.3	Non-parallel version	9
2.4	Parallel version	11
3	Matrix multiplication	14
3.1	Parallel version	18
4	The heat equation	20
4.1	Numeric estimation of second derivatives	26
5	EXERCISES: The heat equation	28

1 Introduction

- For some years now, computers have not become much faster (it seems that in practice there is an upper limit of about 3 GHz).
- Instead, most (all?) computers today come with multiple cores and hence the ability to make parallel computations.
- There are various R packages for parallel computations. The **parallel** package is shipped with the R distribution and this is the package we illustrate here.


2 Parametric bootstrap

2.1 A logistic regression

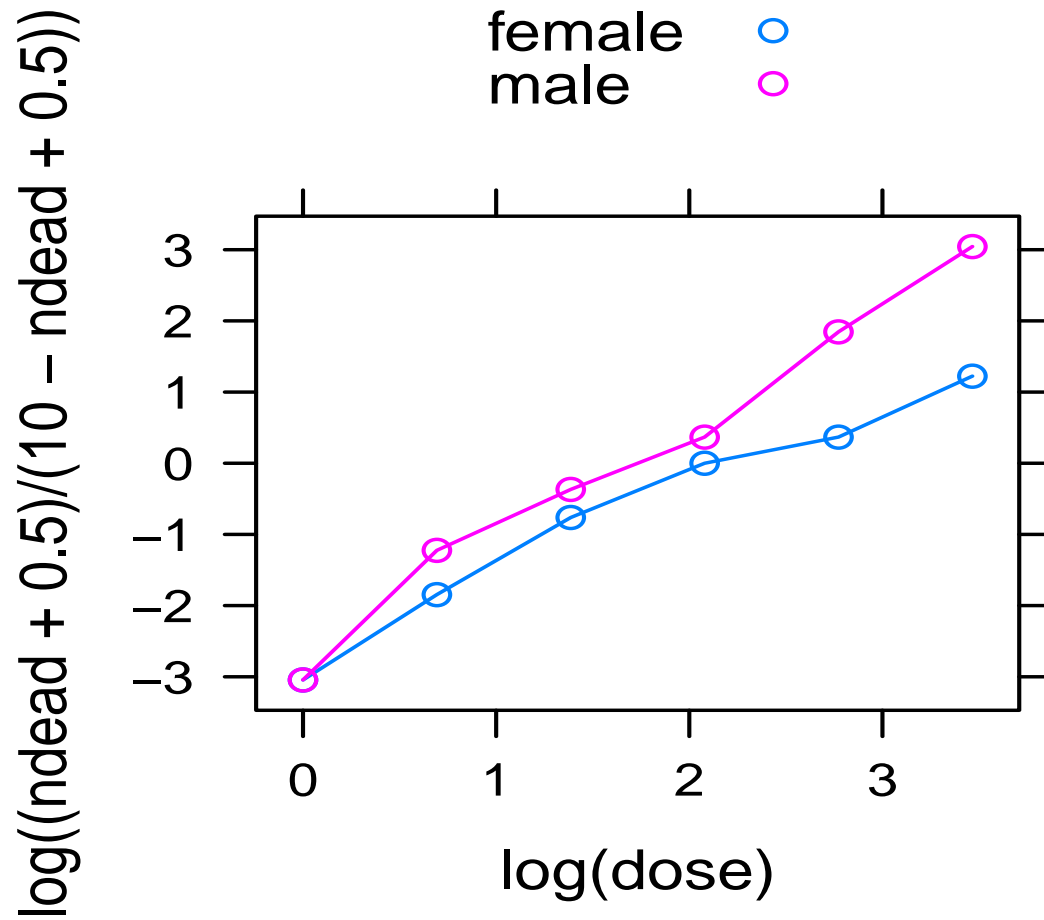
```
R> budworm
```

	sex	dose	ndead	ntotal
1	male	1	0	10
2	male	2	2	10
3	male	4	4	10
4	male	8	6	10
5	male	16	9	10
6	male	32	10	10
7	female	1	0	10
8	female	2	1	10
9	female	4	3	10
10	female	8	5	10
11	female	16	6	10
12	female	32	8	10

```
R> library(lattice)
R> print(xyplot(ndead/10~log(dose), groups=sex, data=budworm,
+             type="b", auto.key=T))
```



```
R> print(xyplot(log((ndead+.5)/(10-ndead+.5))~log(dose), groups=sex, data=budworm,  
+         type="b", auto.key=T))
```



Very vanilla logistic regression. Is there an effect of sex?

```
R> budworm <- transform(budworm, logdose=log(dose))
R> mm1 <- glm(cbind(ndeath, ntotal-ndeath)~sex+logdose,
+           data=budworm, family=binomial(link=logit))
R> mm0 <- update(mm1, .~. - sex)
R> anova(mm1, mm0, test="Chisq")
```

Analysis of Deviance Table

Model 1: cbind(ndeath, ntotal - ndeath) ~ sex + logdose

Model 2: cbind(ndeath, ntotal - ndeath) ~ logdose

	Resid. Df	Resid. Dev	Df	Deviance	Pr(>Chi)
1	9	5.0466			
2	10	8.7883	-1	-3.7417	0.05307 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

2.2 Parametric bootstrap

Alternative to $-2 \log Q$ test: Parametric bootstrap – seems attractive when large sample asymptotics is questionable:

- Compare two models M_1 and M_0 where $M_0 \subset M_1$.
- Fit both models to data; Gives $\hat{\theta}_1$, $\hat{\theta}_0$ and t_{obs} (test statistic)
- Draw B parametric bootstrap sample datasets y^1, \dots, y^B from $p_0(\cdot | \hat{\theta}_0)$. For each simulated dataset y^b , calculate test statistic t^b .
- Evaluate how extreme t_{obs} is in $\{t^1, \dots, t^B\}$.

2.3 Non-parallel version

Need to refit glm to new response variable:

```
R> refit_glm <- function(mm, yy){
+   ff <- update.formula(formula(mm), yy ~ .)
+   ee <- local({ yy=yy; environment() })
+   environment(ff) <- ee
+   cl <- getCall(mm)
+   cl$formula <- ff
+   eval(cl)
+ }
```

Need to do so many times and calculate reference distribution:

```
R> pboot <- function(lg, sm, nsim=10){
+   simdata <- simulate(sm, nsim)
+   ref <- rep.int(NA, nsim)
+   for (ii in 1:nsim){
+     y.new <- simdata[,ii]
+     ref[ii] <- 2*(logLik(refit_glm(lg, y.new)) - logLik(refit_glm(sm, y.new)))
+   }
+   ref
+ }
```

```
R> anova(mm1, mm0, test="Chisq")
```

Analysis of Deviance Table

```
Model 1: cbind(ndeath, ntotal - ndeath) ~ sex + logdose
```

```
Model 2: cbind(ndeath, ntotal - ndeath) ~ logdose
```

	Resid. Df	Resid. Dev	Df	Deviance	Pr(>Chi)
1	9	5.0466			
2	10	8.7883	-1	-3.7417	0.05307 .

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
R> set.seed(123)
```

```
R> Nsim <- 2000
```

```
R> (obsdev <- c(2*(logLik(mm1)-logLik(mm0))))
```

```
[1] 3.741666
```

```
R> system.time({ refdev <- pboot(mm1, mm0, nsim=Nsim) })
```

```
      user  system elapsed
```

```
10.51    0.02    10.55
```

```
R> sum(refdev > obsdev) / length(refdev)
```

```
[1] 0.0505
```

2.4 Parallel version

To do parallel computations, first detect the number of cores available. Next create a handle on the clusters:

```
R> library("parallel")  
R> ## Number of cores:  
R> (nc <- detectCores())  
  
[1] 8  
  
R> ## Create clusters  
R> cl <- makeCluster(rep("localhost", nc))
```

When done with the parallel computations, close the clusters before leaving R:

```
R> ## Remember to shut down clusters before quitting R  
R> stopCluster(cl)
```

```
R> t0 <- proc.time()
R> Nsim <- 2000
R> (nsim2 <- round(Nsim / nc))

[1] 250

R> ## Export global environment to each cluster
R> clusterExport(cl, ls(envir=.GlobalEnv), envir = .GlobalEnv)
R> ## Spec for random number generator
R> clusterSetRNGStream(cl)
R> ## Time to get things up and running
R> proc.time()-t0

  user  system elapsed
0.02   0.03   0.05
```

```
R> t0 <- proc.time()
R> ## Do the simulations
R> xxx <- clusterCall(cl, pboot, mm1, mm0, nsim2)
R> refdev <- c(xxx, recursive=TRUE)
R> ## Time to do the simulations
R> proc.time()-t0

   user  system elapsed 
  0.01    0.00    2.43 

R> (p.PB <- sum(refdev > obsdev) / length(refdev))

[1] 0.0585
```

Notice:

- We get a substantial saving in computing time, but not by a factor equal to the number of cores.
- There is some overhead in setting things up for parallel computations.

3 Matrix multiplication

Multiplication of a $p \times q$ and a $q \times r$ matrix has complexity of the order pqr .

```
R> nr <- 3
R> (A <- matrix(round(rnorm(nr^2),1),nr=nr))

      [,1] [,2] [,3]
[1,] -0.2  0.4 -1.6
[2,]  0.8  0.2  0.5
[3,] -1.2  1.0  0.3

R> (B <- t(A) + 4)

      [,1] [,2] [,3]
[1,]  3.8  4.8  2.8
[2,]  4.4  4.2  5.0
[3,]  2.4  4.5  4.3

R> A %*% B

      [,1] [,2] [,3]
[1,] -2.84 -6.48 -5.44
[2,]  5.12  6.93  5.39
[3,]  0.56 -0.21  2.93
```

Idea: Split A by rows, do multiplications and stack the results:

```
R> A[1,,drop=FALSE] %*% B
```

```
      [,1] [,2] [,3]
[1,] -2.84 -6.48 -5.44
```

```
R> A[2,,drop=FALSE] %*% B
```

```
      [,1] [,2] [,3]
[1,]  5.12  6.93  5.39
```

```
R> A[3,,drop=FALSE] %*% B
```

```
      [,1] [,2] [,3]
[1,]  0.56 -0.21  2.93
```

```
R> rbind(
```

```
+ A[1,,drop=FALSE] %*% B,
```

```
+ A[2,,drop=FALSE] %*% B,
```

```
+ A[3,,drop=FALSE] %*% B )
```

```
      [,1] [,2] [,3]
[1,] -2.84 -6.48 -5.44
[2,]  5.12  6.93  5.39
[3,]  0.56 -0.21  2.93
```

```
R> ## Use 2 clusters
R> cl2 <- cl[1:2]
R> (idx <- splitIndices(nrow(A), length(cl2)))

[[1]]
[1] 1

[[2]]
[1] 2 3

R> (Alist <- lapply(idx, function(ii) A[ii,,drop=FALSE]))

[[1]]
      [,1] [,2] [,3]
[1,] -0.2  0.4 -1.6

[[2]]
      [,1] [,2] [,3]
[1,]  0.8  0.2  0.5
[2,] -1.2  1.0  0.3
```



```
R> ans    <- clusterApply(cl2, Alist, function(aa, BB) aa %*% BB, B)
R> do.call(rbind, ans)
```

```
      [,1] [,2] [,3]
[1,] -2.84 -6.48 -5.44
[2,]  5.12  6.93  5.39
[3,]  0.56 -0.21  2.93
```

```
R> A %*% B
```

```
      [,1] [,2] [,3]
[1,] -2.84 -6.48 -5.44
[2,]  5.12  6.93  5.39
[3,]  0.56 -0.21  2.93
```

3.1 Parallel version

Define parallel matrix multiplication function:

```
R> matprod.par <- function(cl, A, B){
+   if (ncol(A) != nrow(B)) stop("Matrices do not conforme")
+   idx   <- splitIndices(nrow(A), length(cl))
+   Alist <- lapply(idx, function(ii) A[ii,,drop=FALSE])
+   ## ans   <- clusterApply(cl, Alist, function(aa, B) aa %*% B, B)
+   ## Same as above, but faster:
+   ans   <- clusterApply(cl, Alist, get("%*%"), B)
+   do.call(rbind, ans)
+ }
```

Notice:

```
R> get("%*%")

function (x, y)  .Primitive("%*%")
```

```
R> nr <- 5
R> A <- matrix(round(rnorm(nr^2),1),nr=nr)
R> B <- t(A) + 4
R> matprod.par(cl, A, B)
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	-1.65	-2.09	-6.30	-3.94	-5.51
[2,]	-0.49	6.05	-1.91	-3.04	-1.22
[3,]	4.90	7.69	12.41	5.21	10.40
[4,]	13.26	12.56	11.21	17.60	12.50
[5,]	2.09	4.78	6.80	2.90	6.30

```
R> A %*% B
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	-1.65	-2.09	-6.30	-3.94	-5.51
[2,]	-0.49	6.05	-1.91	-3.04	-1.22
[3,]	4.90	7.69	12.41	5.21	10.40
[4,]	13.26	12.56	11.21	17.60	12.50
[5,]	2.09	4.78	6.80	2.90	6.30

4 The heat equation

For a function $u(x, y, t)$ of spatial variables (x, y) and time variable t , the heat equation is


$$\frac{\partial u}{\partial t} - \alpha \left(\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} \right) = 0, \quad \alpha > 0$$

This equation describes how the temperature $u(x, y, t)$ at location (x, y) changes over time t as heat spreads in space.

To solve this equation we must specify the initial conditions $u(x, y, 0)$ and a set of boundary conditions.

Suppose we throw a heated horse shoe (with temperature 1) on a cold plate (with temperature 0):

```
R> nr <- nc <- 30  
R> mm <- matrix(0,nrow=nr, ncol=nc)  
R> mm[7,9:15] <- mm[8:16,16] <- mm[8:16,8] <- 1  
R> image(mm)
```



The temperature at the boundaries (outermost rows and columns) is assumed to remain constantly equal to zero. (Called a Dirichlet condition).

On a regular grid, this equation can be solved numerically as

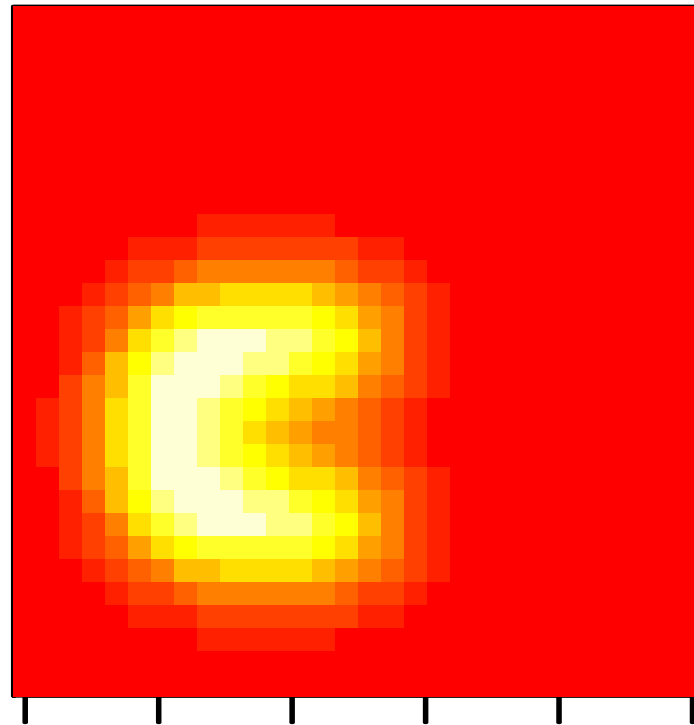
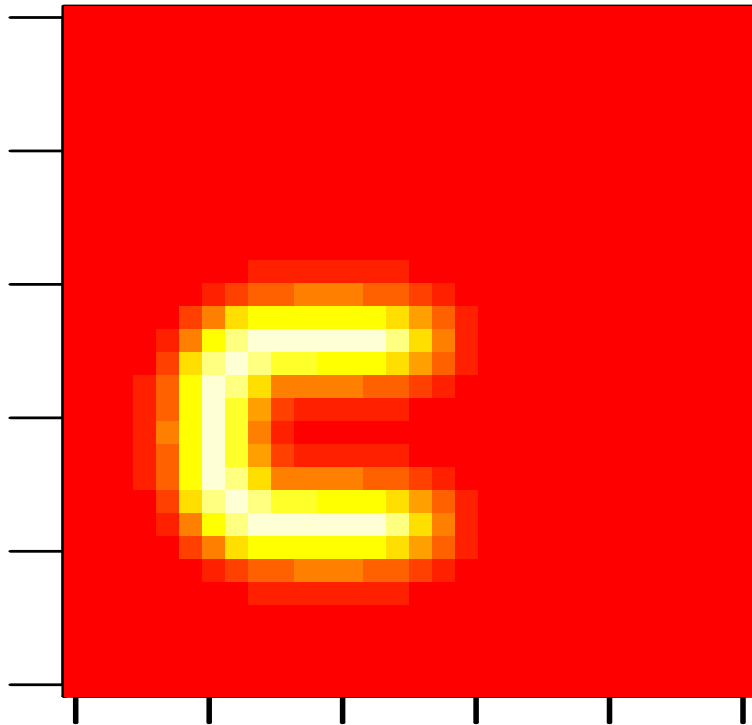
$$\begin{aligned} u(x, y, t + 1) = & u(x, y, t) + \\ & \alpha(u(x + 1, y, t) + u(x - 1, y, t) - 2u(x, y, t)) + \\ & \alpha(u(x, y + 1, t) + u(x, y - 1, t) - 2u(x, y, t)) \end{aligned}$$

```

R> ## Update in one time step
R> heat <- function(mm, cx=.1, cy=.1, nr=nrow(mm), nc=ncol(mm)){
+   mm2 <- mm
+   for (ii in 2:(nr-1)){
+     for (jj in 2:(nc-1)){
+       mm2[ii,jj] <- mm[ii,jj] +
+         cy * (mm[ii-1,jj]+mm[ii+1,jj]-2*mm[ii,jj]) +
+         cx * (mm[ii,jj-1]+mm[ii,jj+1]-2*mm[ii,jj])
+     }
+   }
+   mm2
+ }
R> library(compiler) ## Byte compiling gives a factor 5 in speed
R> heat.cmp <- cmpfun(heat)
R> ## Iterate over time
R> heat.iter <- function(mm, n=1, cx=.1, cy=.1){
+   for (ii in 1:n){
+     mm <- heat.cmp(mm, cx, cy)
+   }
+   mm
+ }

```

```
R> mm.10 <- heat.iter(mm,n=10)
R> mm.30 <- heat.iter(mm.10,n=20)
R> image(mm.10)
R> image(mm.30,yaxt='n')
```



4.1 Numeric estimation of second derivatives

Suppose that we want to estimate the second derivative of a function f at a point x and that we know the values of the function at three points

$$(x - h, f(x - h)), \quad (x, f(x)), \quad (x + h, f(x + h))$$

Since $f''(x)$ is the derivative of the function f' at the point x we can estimate $f''(x)$ by

$$f''(x) \approx \frac{f'(x + \frac{h}{2}) - f'(x - \frac{h}{2})}{h}$$

We can estimate $f'(x - h/2)$ and $f'(x + h/2)$ by

$$\begin{aligned} f'(x - \frac{h}{2}) &\approx \frac{f(x) - f(x - h)}{h} \\ f'(x + \frac{h}{2}) &\approx \frac{f(x + h) - f(x)}{h} \end{aligned}$$

Putting this together gives

$$\begin{aligned} f''(x) &\approx \frac{\frac{f(x+h)-f(x)}{h} - \frac{f(x)-f(x-h)}{h}}{h} \\ &= \frac{f(x+h) - 2f(x) + f(x-h)}{h^2} \end{aligned}$$

Partial derivatives can be estimated the same way:

$$\begin{aligned} \frac{\partial^2}{\partial x^2} f(x, y) &\approx \frac{f(x+h, y) - 2f(x, y) + f(x-h, y)}{h^2} \\ \frac{\partial^2}{\partial y^2} f(x, y) &\approx \frac{f(x, y+h) - 2f(x, y) + f(x, y-h)}{h^2} \end{aligned}$$

5 EXERCISES: The heat equation

The update over time for the heat equation is an obvious candidate for parallel computing:

1. Divide the grid into regions – for example by splitting by the rows.
2. Update the regions.
3. Put the updated regions together and make the necessary modifications on the boundary between the regions.

Exercise:

1. Implement a parallel version of the algorithm for updating the heat equation.